

Topics

1. GCC Toolchain
2. Shell Programming
3. File Operations
4. Advanced File Operations
5. Linux Kernel Compilation and build procedure
6. Process Management
7. Inter-Process Communication

GNU Toolchain

- GCC (GNU Compiler Collection)
- GNU Makefile
- GDB (GNU Debugger)
- Libraries:
- Procedure for creation of Static and Dynamic Libraries.
- Memory Layout of C program.

File Operations

- Linux File Structure
- Difference between System call and Standard Libraries
- /Proc and /Sys file Systems

Linux Kernel Compilation and build procedure

- Get Latest source code from Kernel
- Extrac source file
- Configure Kernel
- Compile Kernel
- Install Kernel
- Create on initrd image
- Update grub
- Reboot


Process Management

Processes

- Process Concept
 - Process ID & State
 - Process Tree
 - Process Control Block
 - Context Switching
 - Queues
 - Process Termination
 - Viewing Processes
- Process Scheduling
 - Linux Scheduling
 - RT Issues
- Process Creation
 - Replacing Process Image
 - Replication of Processes
 - Waiting for Processes

Process Concept

- Like files, a process is a fundamental abstraction in Unix/Linux
 - An executing instance of a program
- A process is an “an address space with one or more threads executing within that address space, and the required system resources for those threads.”
- The Linux kernel, supporting both preemptive multitasking and virtual memory, provides a process both a virtualized processor and a virtualized view of memory
- Each process consists of one or more *threads* of execution
- A thread is the unit of activity within a process, the abstraction responsible for executing code
- Each thread has
 - an id (*pid*)
 - a stack
 - state
 - program counter



Process Concept

- Parent processes create children processes, which, in turn create other processes, forming a tree of processes (*process tree*)
 - init process (pid=1) is the first created after booting.
- Resource sharing between parent and child processes
 - Parent and children share all resources
 - Children share subset of parent's resources (process share groups)
 - Parent and child share no resources
 - Resource sharing in Linux
 - ▶ File structure, resource limits, pipes, FIFOs etc
- Execution
 - Parent and children execute concurrently
 - Parent can wait until children terminate

Process ID (pid)

- Each process has a unique identifier, the *process ID* (maximum 32768)
- The process ID is represented by the `pid_t` type, defined in `<sys/types.h>`
- The `getpid()` system call returns the process ID of the invoking process
- The `getppid()` system call returns the ID of the parent of the invoking process

```
#include <sys/types.h>
```

```
#include <unistd.h>
```

```
#include <stdio.h>
```

```
int main() {
```

```
    printf ("My pid=%d\n", getpid ( ));
```

```
    printf ("Parent's pid=%d\n", getppid ( ));
```

```
    return (0);
```

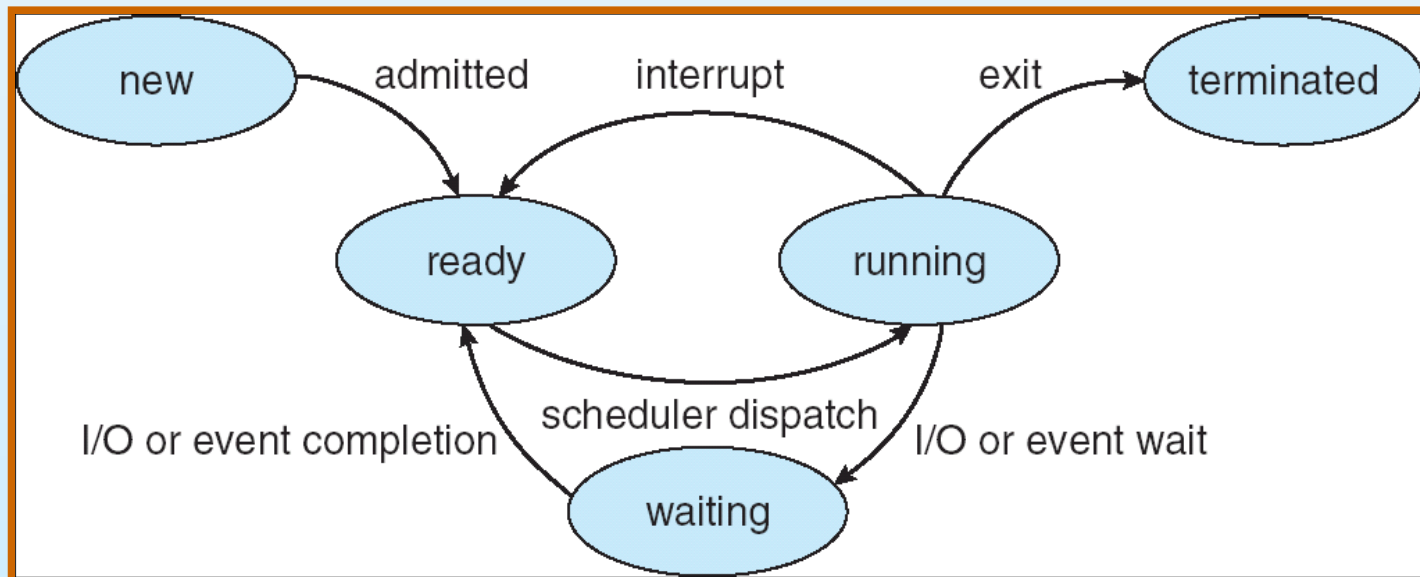
```
}
```

My pid=6811

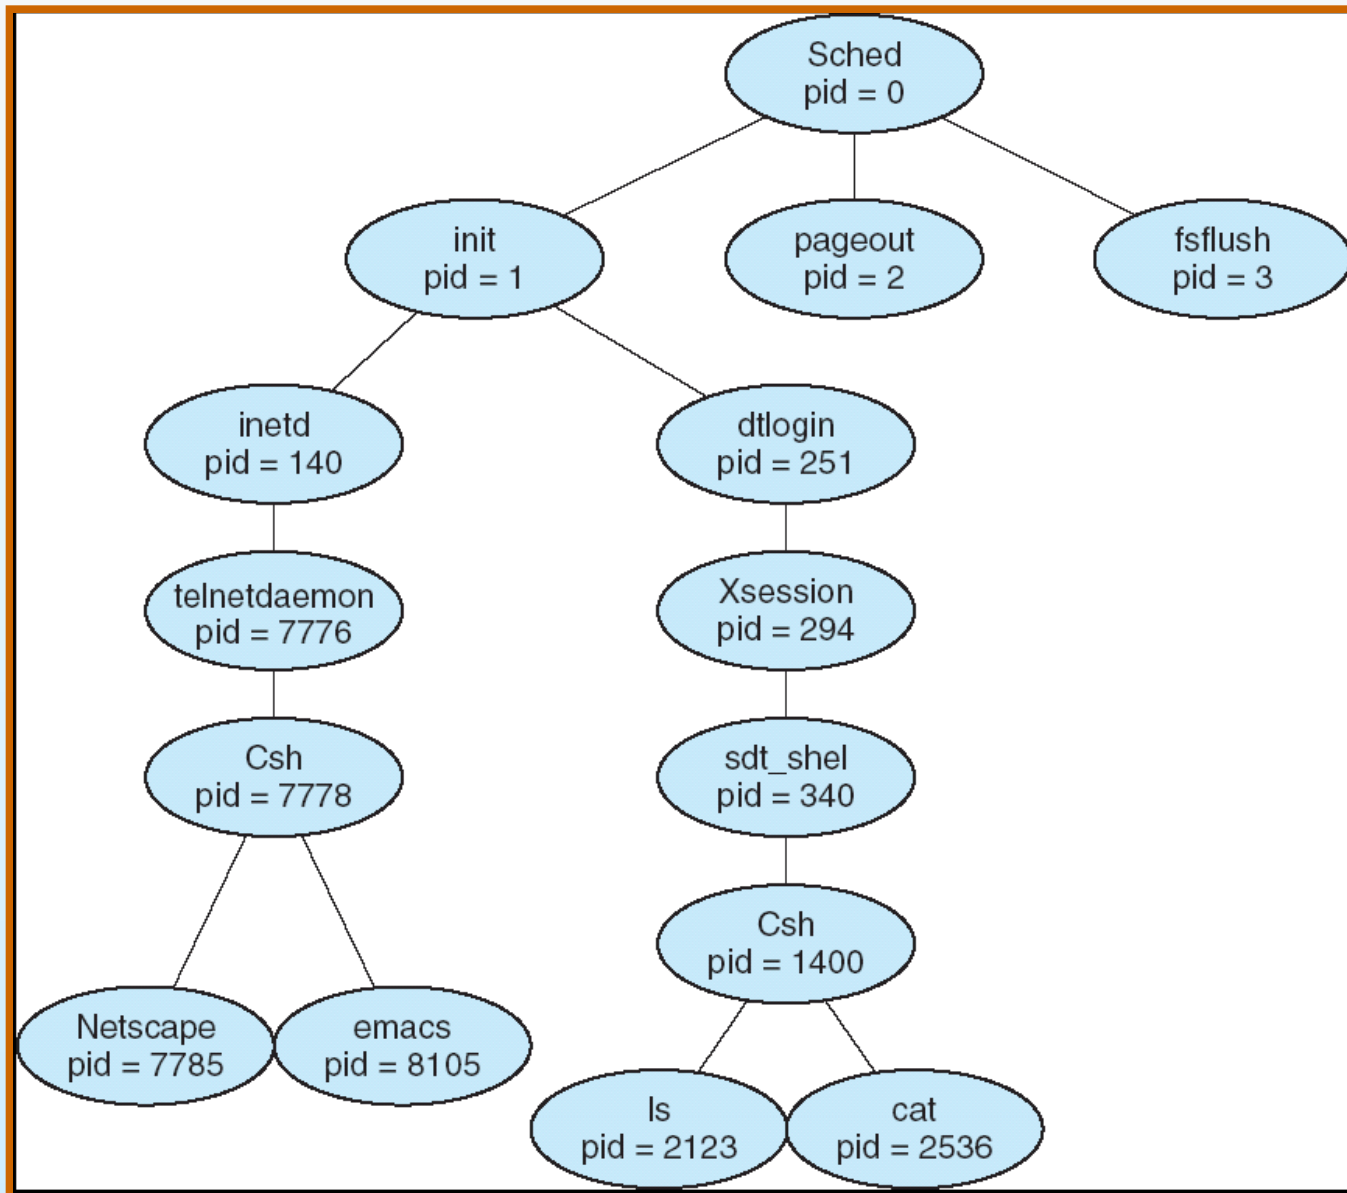
Parents pid=6723

Process State

- As a process executes, it changes *state*
 - new**: The process is being created
 - running**: Being executed
 - waiting**: The process is waiting for some event to occur
 - ready**: The process is waiting to be assigned to a processor
 - terminated**: The process has finished execution
- State values: TASK_RUNNING, TASK_INTERRUPTIBLE, TASK_UNINTERRUPTIBLE, TASK_STOPPED, TASK_ZOMBIE



A tree of processes



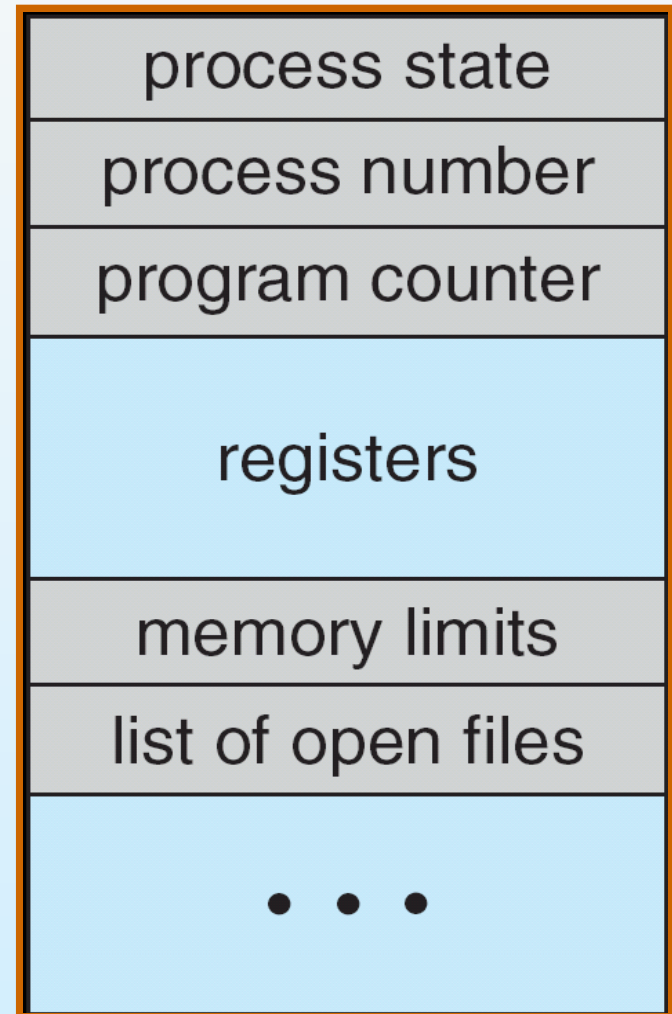
Init Process

- The first process that the kernel executes after booting the system, called the *init process*, has the pid 1
- The init process handles
 - The remainder of the boot process
 - Initializing the system
 - Starting various services
 - Launching a login program
- The Linux kernel tries four executables, in the following order:
 - */sbin/init*: The preferred and most likely location for the init process.
 - */etc/init*: Another likely location for the init process.
 - */bin/init*: A possible location for the init process.
 - */bin/sh*: The Bourne shell, if it fails to find an init process

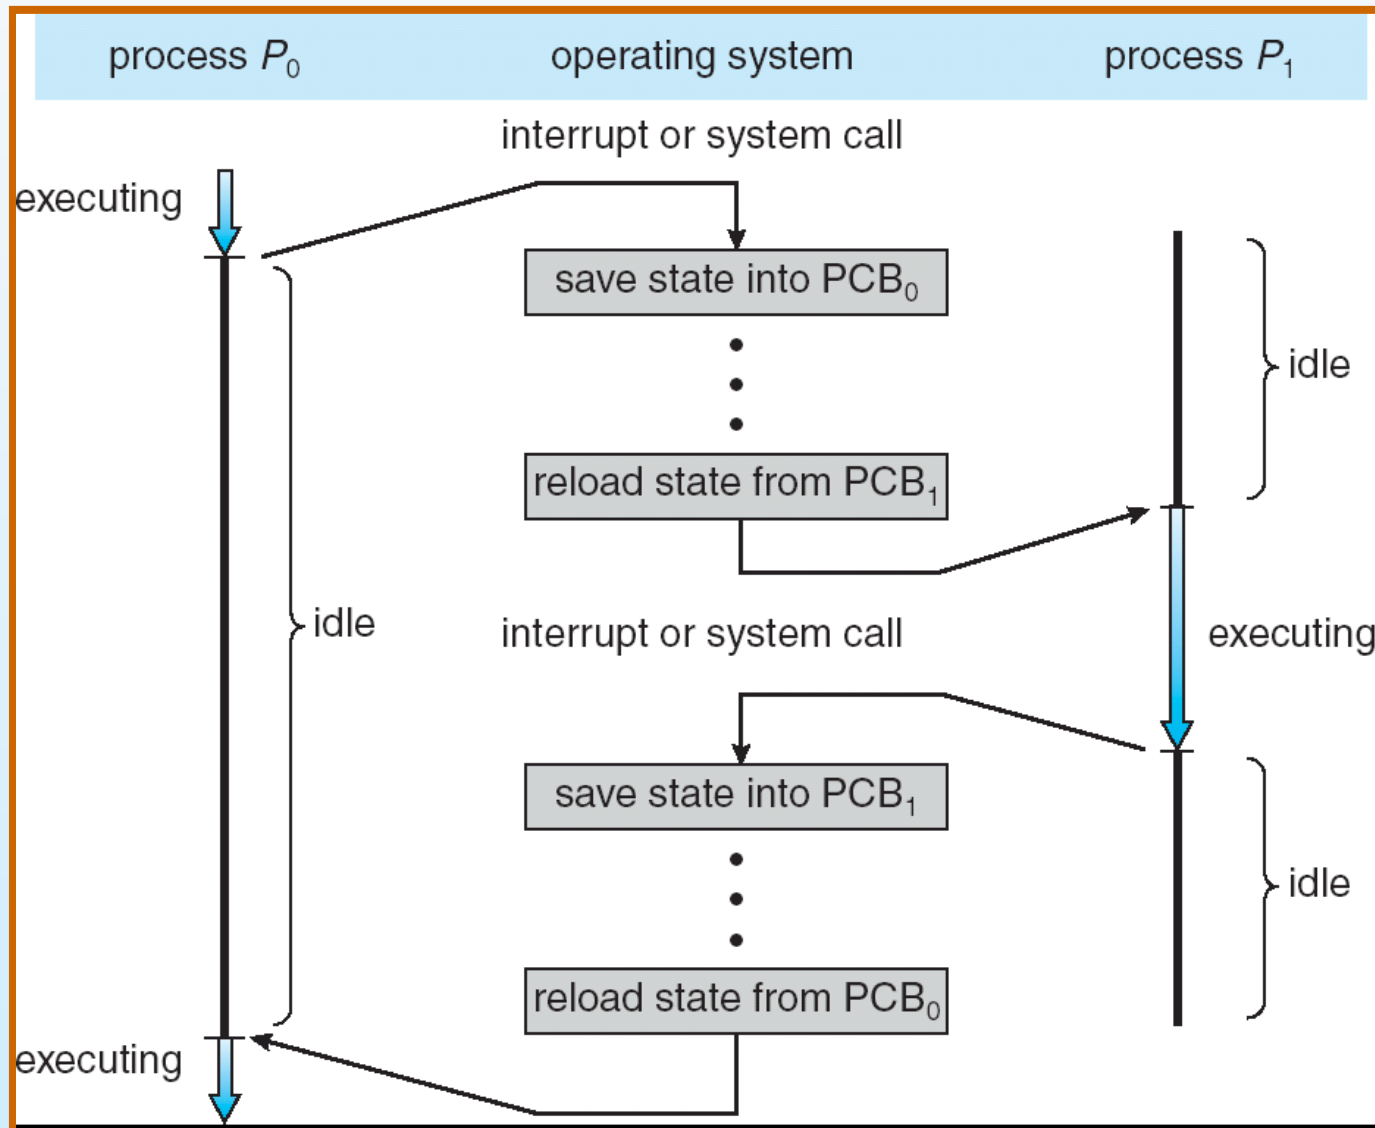
Process Control Block (PCB)

Information associated with each process stored in a block of memory known as PCB or Process Descriptor

- ▢ Process state
- ▢ Program counter
- ▢ CPU registers
- ▢ CPU scheduling information
- ▢ Memory-management information
- ▢ Accounting information
- ▢ I/O status information



CPU Switch From Process to Process



Process Termination

- ❑ Process executes last statement (**exit 0** for successful exit, **exit 1**, or >0 for exit with error condition) to inform the operating system to delete it
 - ❑ Process' resources are deallocated by operating system
- ❑ Parent may terminate execution of children processes (**abort**)
- ❑ Parents may wait (via **wait**) for a child process to terminate
 - ❑ If a child process terminates before the parent does wait, Linux does not delete it fully but keeps the exit information for the parent (**zombie**)
- ❑ If a parent process exits
 - ❑ Some operating system do not allow child to continue if its parent terminates
 - ▶ All children terminated - *cascading termination*
 - ❑ In Linux, if a parent terminates before a child, the child is re-parented to another process in the group or to the init process
- ❑ The library call `exit()` is a wrapper over the kernel syscall `_exit()`. `exit()` flushes pending I/O, closes file descriptors and does other cleanup (memory, semaphores, etc) before calling `_exit()`

Viewing Processes

- Linux Process Table
 - a data structure describing all of the processes that are currently loaded
- Viewing processes
 - The **ps** command shows the processes in the system or belonging to a user

\$ ps -af

```
UID PID PPID C STIME TTY  TIME  CMD
root 433  425  0 18:12  tty1 00:00:00 [bash]
```

- Process priority

\$ ps -l

```
F  S UID  PID  PPID C PRI NI  SZ  WCHAN TTY  TIME  CMD
000 S 500 1362 1262 2 80   0 789  schedu pts/1 00:00:00 oclock
```


Scheduling Methods

- ▢ First-Come, First-Served (FCFS)
 - ▢ Non-preemptive; convoy effect
- ▢ Shortest-Job-First (SJF) Scheduling
 - ▢ Non-preemptive & Preemptive (SRTF) – optimal
 - ▢ Issue – weighted average prediction of length of next CPU burst
- ▢ Priority Scheduling
 - ▢ Starvation – increase priority of the process with time (Aging)
 - ▢ For equi-priority processes
 - ▶ FIFO
 - ▶ RR (Round Robin) – time slice (10-100 ms); context switch – 100μs
- ▢ Multilevel queues – (foreground/background) with own scheduling policy
- ▢ Multiple-Processor Scheduling
 - ▢ Asymmetric Multiprocessing – one CPU (master) handles scheduling
 - ▢ Symmetric Multiprocessing – each processor is self-scheduling
 - ▶ Processor affinity to improve cache performance
 - ▶ Load Balancing by Push and Pull migration

Linux Scheduling

- Preemptive, priority-based scheduling
- Two algorithms: time-sharing and real-time
- Time-sharing (nice values range from 100 to 140)
 - Dynamic priorities
 - ▶ I/O bound jobs have higher priorities
 - ▶ Aging adds priority
- Real-time (Real-time priorities vary from 0 to 99)
 - Soft real-time
 - Posix.1b compliant – (static priorities)
 - ▶ FCFS and RR

RT issues in Linux

- A real time system can be defined as a *"system capable of guaranteeing timing requirements of the processes under its control"*
 - Fast – low latency (quick response to external, asynchronous events)
 - Predictable – deterministic completion time of tasks with certainty
- Linux is a standard time-sharing operating system with good average performance and highly sophisticated services but lacks real time support
 - Changes in the interrupt handling
 - Scheduling policies
- Two Approaches to make Real Time OS
 - Develop new Linux like RTOS – QNX Neutrino, Vx Works, Lynx, µSoft CE
 - Plug-in a layer (Real-time kernel) underneath the Linux Kernel to handle interrupts and do real-time scheduling
- A real time kernel intercepts the interrupts, not allowing Linux to disable them, and decides what to despatch – schedule tasks or send to Linux
 - **RTLinux** was developed at the New Mexico Institute of Technology by Michael Barabanov and Professor Victor Yodaiken
 - **Real-Time Application Interface (RTAI)** was developed at the Dipartimento di Ingegneria Aerospaziale, Politecnico di Milano (DIAPM) by Professor Paolo Mantegazza

Starting new Processes

□ Starting New Processes

// system1.c to create a new process by
//using the **system** library function

```
#include <stdlib.h>
```

```
#include <stdio.h>
```

```
int main()
```

```
{
```

```
    printf("Running ps with system\n");
```

```
    system("ps -ax");
```

```
    printf("Done.\n");
```

```
    exit(0);
```

```
}
```

// to run in background system2.c

// replace system("ps -ax") with

```
system("ps -ax" &);
```

```
$ ./system1
```

Running ps with system

```
PID TTY STAT TIME COMMAND
```

```
- - -
```

```
1480 pts/1 S 0:00 ./system1
```

```
1481 pts/1 R 0:00 ps -ax
```

Done.

```
$ ./system2
```

Running ps with system

```
PID TTY STAT TIME COMMAND
```

```
1 ? S 0:05 init
```

```
2 ? SW 0:00 [keventd]
```

```
...
```

Done.

```
$ 1483 ? S 0:00 kdeinit: klipper
```

```
1484 pts/1 R 0:00 ps -ax
```

Replacing a Process image

- **exec** function replaces the current process with a new process specified by the path or file argument

```
int execl (const char *path, const char *arg0, ..., (char *)0);
```

```
int execlp (const char *file, const char *arg0, ..., (char *)0);
```

```
int execlp (const char *path, const char *arg0, ..., (char *)0, char *const envp[]);
```

```
int execv (const char *path, char *const argv[]); //basic syscall
```

```
int execvp (const char *file, char *const argv[]);
```

```
int execve (const char *path, char *const argv[], char *const envp[]);
```

“l” indicates that the arguments are provided in a null terminated list; “v” in an array (vector);

“p” indicates the full PATH must be searched for the file;

- “e” indicates a new environment is also supplied for the new process

- **ret = execl(“/bin/ps”, “ps”, “-ax”, 0);** /* assumes ps is in /bin */

- replaces the current process image by loading the program pointed at by path

- **ret = execlp(“ps”, “ps”, “-ax”, 0);** /* assumes /bin is in PATH */

- To use the “v” or array option

```
const char *args[ ] = { “ps”, “-ax”, NULL };
```

- **ret = execv (“/bin/ps”, args);** or **ret = execvp (“ps”, args);**

exec call

- A successful invocation of exec call does not return; it ends by jumping to the entry point of the new program, and the just-executed code no longer exists in the process' address space
- On error execl() returns -1, and sets errno to indicate the problem (examples of errno values: EACCESS, ENOEXEC, ENOMEM, etc)
Note: errno variable is defined in <errno.h> include file
- On successful exec call
 - some properties of process are same: pid, priority, owning user and group
 - some properties change: signals, memory locks, statistics
 - open files are retained; generally these are closed before the exec call

Replacing a Process image

▮ pexec1p

```
#include <unistd.h>
```

```
#include <stdio.h>
```

```
int main()
```

```
{
```

```
    int rtn;
```

```
    printf("Running ps with execlp\n");
```

```
    rtn = execlp("ps", "ps", "-ax", 0);
```

```
    if (ret == -1)
```

```
        perror ("execl");
```

```
    printf("Done.\n");
```

```
    exit(0);
```

```
}
```

\$./pexec1p

Running ps with execlp

```
PID TTY STAT TIME
      COMMAND
```

```
1 ? S 0:05 init
```

```
2 ? SW 0:00 [keventd]
```

```
...
```

```
1262 pts/1 S 0:00 /bin/bash
```

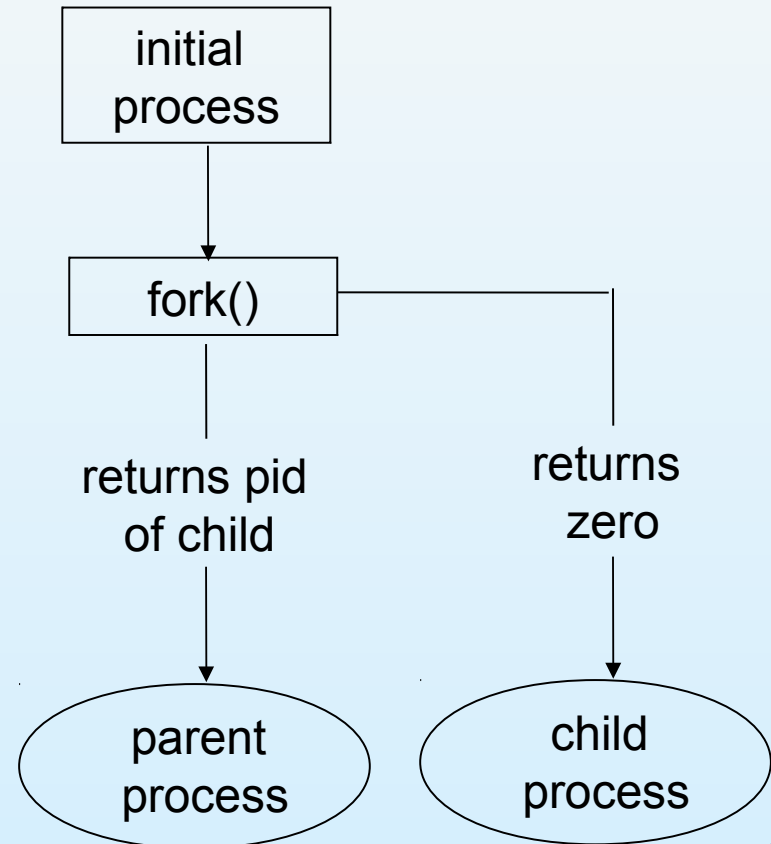
```
1273 pts/2 S 0:00 su -
```

```
1274 pts/2 S 0:00 -bash
```

Duplicating a Process Image

- We can create a new process by calling **fork**. This system call duplicates the current process (creates a new entry in the process table with same attributes as the current process)
- Both processes continue from next instruction

```
pid_t new_pid;  
- - -  
new_pid = fork();  
switch(new_pid) {  
    case -1 : /* Error */  
        break;  
    case 0 : /* We are child */  
        - - -  
        break;  
    default : /* We are parent */  
        - - -  
        break;  
}
```



fork

```
#include <sys/types.h>
#include <unistd.h>
#include <stdio.h>
int main()
{
    pid_t new_pid; char * msg; int n;
    printf ("fork program starting\n");
    new_pid = fork();
    switch(new_pid) {
        case -1: //fail
            perror ("fork failed"); exit(1);
        case 0: //child
            msg = "child"; n=5; break;
        default: //parent
            msg = "parent"; n=3; break;
    }
```

```
for (; n>0; n--) {
    puts(msg); sleep(1)
}
exit (0);
}
```

❏ \$./fork1

fork program starting

This is the parent

This is the child

This is the parent

This is the child

This is the parent

This is the child

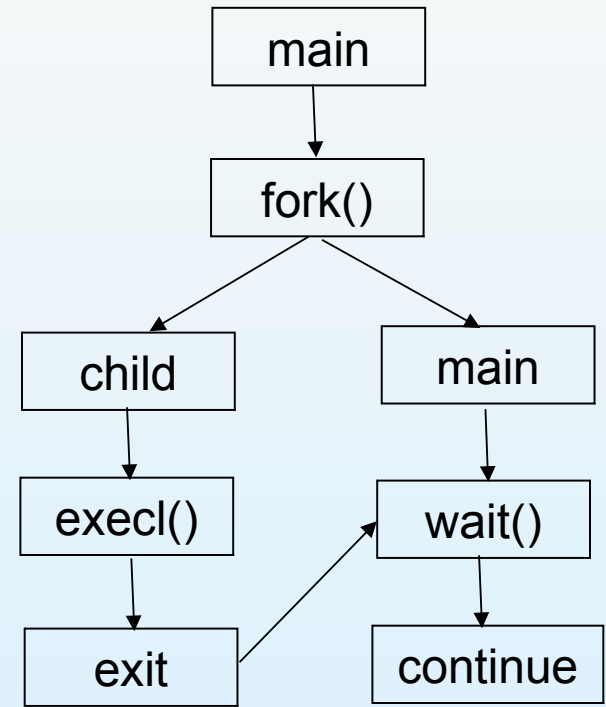
\$ This is the child

This is the child

fork and exec

```
#include <sys/wait.h>

int main()
{
    pid_t pid;
    /* fork another process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        exit(-1);
    }
    else if (pid == 0) { /* child process */
        execl ("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait (NULL);
        printf ("Child Complete");
        exit(0);
    }
}
```



fork call

□ The successful fork() call

- The fork() call makes a copy of the parent process structure for the child
 - ▶ Address space, resource limits, umask, controlling terminal, directory structure, current working directory, file pointers etc
- The following will be different
 - ▶ PID, PPID, resource utilizations (child set to 0), signals etc

□ On failure

- fork() returns – 1, and error set in errno (EAGAIN, ENOMEM)

Waiting for a Process – wait()

- Parent process can wait for the child to finish by calling
pid_t wait (int *stat_val);
- The call returns PID & exit status of the child process in stat_val
- Need macros to interpret
 - WIFEXITED(stat_val) – Nonzero if the child is terminated normally
 - WEXITSTATUS(stat_val) – child exit code If WIFEXITED is nonzero
 - WIFSIGNALED(stat_val) – Nonzero if child terminated on uncaught signal
 - WTERMSIG(stat_val) – signal number if WIFSIGNALED is nonzero
- To wait for a specific process
pid_t waitpid (pid_t pid, int *status, int options);
 - Options WNOHANG – Do not block

Waiting for a Process – wait()

```
#include <sys/wait.h>

int main() {
    pid_t pid;
    ...

    pid = fork();
    switch(pid)
    {
    case -1:
        ...
    case 0:
        printf("This is the child\n");
        wait(10);
        exit(0);
    default:
        ...
    }

    if (pid != 0) {    //Parent
        int stat_val;
        pid_t child_pid;
        child_pid = wait(&stat_val);
        printf("Child finished: PID = %d\n", child_pid);
        if(WIFEXITED(stat_val))
            printf("Child exited with code %d\n",
                WEXITSTATUS(stat_val));
        else
            printf("Child terminated abnormally\n");
    }
    exit(exit_code);
}
```

```
$ ./wait
This is the child
Child finished: PID = 1582
Child exited with code 37
$
```

Process Priority

- Unix has historically called process priorities *nice values*
 - Legal nice values range from -20 to 19 inclusive, (default value of 0) (the lower a process' nice value, the higher its priority)
 - Linux provides system calls for retrieving and setting a process' nice value
- **int nice (int inc);**
 - If inc = 0, nice returns current value
 - For inc > 0, nice increments the nice value by inc & returns the new value
- **getpriority(), setpriority(), renice()**
 - Get and set priority for individual process, group or user

```
#include <unistd.h>
#include <stdio.h>
int main()
{
    int ret, val;
    val = nice (0);
    printf ("current nice
value is %d\n", val);
```

```
ret = nice (10); /* increase nice by 10 */
if (ret == -1)
    perror ("nice");
else
    printf ("nice value is now %d\n",
ret);
return 0;
}
```

current nice value is 0
Nice value is now 10

Signals

- *Signals* are software interrupts for handling asynchronous events
 - External – eg. the interrupt character (Ctrl-C)
 - Internal – as when the process divides by zero
 - A process can also send a signal (“raise”) to another process
- Signal life cycle
 - A signal is “*raised*”
 - Kernel *stores* and *delivers the signal*
 - The process *handles* the signal
- Signal handling
 - Ignore the signal – 2 cannot be ignored – SIGKILL & SIGSTOP
 - Catch and handle the signal by registered functions (signal handlers)
 - SIGINT and SIGTERM are two commonly caught signals
 - Default action – terminate the process (result in core dump)

Signals

- Signals have identifiers defined in `<signal.h>`
(They also have integers)
- Processes can handle a signal they catch by defining a signal function to handle them
(void) signal(SIGINT, ouch);
- A successful call to `signal()` removes the current action taken on receipt of the signal `SIGINT`, and instead handles the signal with the signal handler specified by the handler function `ouch(int)`
- To reset the signal to the default behavior or to ignore by
 - SIG_DFL** – default handler
 - SIG_IGN** – ignore the signal

Signal Name	Description
SIGABORT	Process abort
SIGALRM	Alarm clock
SIGFPE	Floating-pt exception
SIGHUP	Hangup
SIGILL	Illegal Instruction
SIGINT	Terminal interrupt
SIGKILL	Kill
SIGPIPE	Write on pipe, with no reader
SIGQUIT	Terminal quit
SIGSEGV	Invalid memory segment access
SIGTERM	Termination (ctrl-c)

Signals

ctrlc1.c

```
#include <signal.h>
```

```
#include <stdio.h>
```

```
#include <unistd.h>
```

```
void ouch(int sig)
```

```
{
```

```
    printf("OUCH! - I got signal %d\n", sig);
```

```
    (void) signal(SIGINT, SIG_DFL);
```

```
}
```

```
int main()
```

```
{
```

```
    (void) signal(SIGINT, ouch);
```

```
    while(1) {
```

```
        printf("Hello World!\n");
```

```
        sleep(1);
```

```
    }
```

```
}
```

```
$ ./ctrlc1
```

```
Hello World!
```

```
Hello World!
```

```
Hello World!
```

```
Hello World!
```

```
^C
```

```
OUCH! - I got signal 2
```

```
Hello World!
```

```
Hello World!
```

```
Hello World!
```

```
Hello World!
```

```
^C
```

```
$
```

Signals

- A child process does not inherit the signal handlers from the parent
- They are set to default handlers
- A process may send a signal to another process with same user ID, including itself, by calling kill.
#include <sys/types.h>
#include <signal.h>
int kill(pid_t pid, int sig);
- The **kill** function sends the specified signal, **sig**, to the process whose identifier is given by pid.
- **int raise (int signo)** function sends the signal **sigio** to itself
- Signals provide us with a useful alarm clock facility. The alarm function call can be used by a process to schedule a SIGALRM signal at some time in future.
#include <unistd.h>
unsigned int alarm(unsigned int seconds);
- A robust signal interface
 - **sigaction()** and **sigqueue()**

An Alarm Clock (alarm.c)

```
static int alarm_fired = 0;
void ding(int sig)
{
    alarm_fired = 1;
}
int main() {
    pid_t pid;
    printf("alarm application
        starting\n");
    pid = fork();
    switch(pid) {
    case -1: /* Failure */
        perror("fork failed");
        exit(1);
    case 0: /* child */
        sleep(5);
        kill(getppid(), SIGALRM);
        exit(0);
    }
```

```
/* This is the parent process */
printf("waiting for alarm to go
    off\n");
(void) signal(SIGALRM, ding);
pause();
if (alarm_fired) printf("Ding!\n");
printf("done\n");
exit(0);
}
```

```
$ ./alarm
alarm application starting
waiting for alarm to go off
<5 second pause>
Ding!
done
$
```

A process is a program in execution. It consists of the executing program code, a set of resources such as open files, internal kernel data, an address space, one or more threads of execution and a data section containing global variables.

Fork creates a new process which is a copy of the calling process. That means that it copies the callers memory (**code, globals, heap and stack**), **registers and file descriptors**.

Difference between process and thread:

Process	Thread
Process is an execution of a program. Program by itself is not a process.	Thread is a smallest unit of execution (Light weight process). Thread is a sequence of control within a process.
Whenever you are creating a process text segment, data segment, stack segments are created.	Whenever you are creating a thread it will create only stack segment and share text and data segment. i.e., why we are saying thread is smallest unit of execution.
Process is a parallel execution.	Thread is a serial execution.

Drawbacks of Thread

Debugging a multithreaded program is much harder than debugging a single-threaded one, because the interactions between the threads are very hard to control.

Writing multithreaded programs requires very careful design.

A program that splits a large calculation into two and runs the two parts as different threads will not necessarily run more quickly on a single processor machine, as there are only so many CPU cycles to be had, though if nothing else